

## program1-chronic


March 29, 2024

```
[77]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
sns.set_theme(color_codes=True)
```

```
[78]: df=pd.read_csv('ChronicDataset.csv')
```


```
[79]: sns.countplot(data=df, hue="Class", x="Htn")
```

```
[79]: <Axes: xlabel='Htn', ylabel='count'>
```




```
[80]: sns.countplot(data=df , hue="Class", x="Rbc")
```

```
[80]: <Axes: xlabel='Rbc', ylabel='count'>
```



```
[81]: sns.histplot(data=df, hue="Class", x="Bp" ,multiple="stack")
```

```
[81]: <Axes: xlabel='Bp', ylabel='Count'>
```



## Data Pre-Processing

```
[82]: df.isnull().sum()
```

```
[82]: Bp      0
      Sg      0
      Al      0
      Su      0
      Rbc     0
      Bu      0
      Sc      0
      Sod     0
      Pot     0
      Hemo    0
      Wbcc    0
      Rbcc    0
      Htn     0
      Class   0
      dtype: int64
```

```
[83]: df_copy=df.copy(deep=True)
```


```
[84]: df_copy[['Bp', 'Sg', 'Bu', 'Sc', 'Sod', 'Pot', 'Hemo', 'Wbcc', 'Rbcc']] =  
      ↪df_copy[['Bp', 'Sg', 'Bu', 'Sc', 'Sod', 'Pot', 'Hemo', 'Wbcc', 'Rbcc']].  
      ↪replace(0, np.nan)
```

```
[85]: print(df_copy.isnull().sum())
```

```
Bp      0  
Sg      0  
Al      0  
Su      0  
Rbc     0  
Bu      0  
Sc      0  
Sod     0  
Pot     0  
Hemo    0  
Wbcc    0  
Rbcc    0  
Htn     0  
Class   0  
dtype: int64
```

```
[86]: sns.countplot(data=df ,x='Class',hue='Class')  
      print(df.Class.value_counts())
```

```
Class  
1      250  
0      150  
Name: count, dtype: int64
```




Balancing the data

```
[87]: from sklearn.utils import resample
df_maj=df[(df['Class']==1)]
df_min=df[(df['Class']==0)]
df_min_up=resample(df_min,
                    n_samples=250,
                    random_state=0)
df2=pd.concat([df_min_up,df_maj])
```


```
[88]: sns.countplot(data=df2, x="Class", hue="Class")
print(df2.Class.value_counts())
```

```
Class
0    250
1    250
Name: count, dtype: int64
```




```
[89]: sns.boxplot(x=df2["Bp"])
```

```
[89]: <Axes: xlabel='Bp'>
```



```
[90]: sns.boxplot(x=df2["Sg"])
```


```
[90]: <Axes: xlabel='Sg'>
```



```
[91]: sns.boxplot(x=df2["Bu"])
```


```
[91]: <Axes: xlabel='Bu'>
```






```
[92]: sns.boxplot(x=df2["Sc"])
```

```
[92]: <Axes: xlabel='Sc'>
```




```
[93]: sns.boxplot(x=df2["Sod"])
```

```
[93]: <Axes: xlabel='Sod'>
```




```
[94]: sns.boxplot(x=df2["Pot"])
```

```
[94]: <Axes: xlabel='Pot'>
```




```
[95]: sns.boxplot(x=df2["Hemo"])
```

```
[95]: <Axes: xlabel='Hemo'>
```




```
[96]: sns.boxplot(x=df2["Wbcc"])
```

```
[96]: <Axes: xlabel='Wbcc'>
```



```
[97]: sns.boxplot(x=df2["Rbcc"])
```

```
[97]: <Axes: xlabel='Rbcc'>
```




```
[98]: import scipy.stats as stats
      z = np.abs(stats.zscore(df2))
      data_clean = df2[(z<3).all(axis = 1)]
      data_clean.shape
```

```
[98]: (420, 14)
```

```
[99]: sns.heatmap(data_clean.corr(), fmt='.2g')
```


```
[99]: <Axes: >
```



```
[100]: #Rbc attribute is irrlevant, so we have to remove it
data_clean2 = data_clean.drop(columns=['Rbc'])
```

```
[101]: corr = data_clean2[data_clean2.columns[1:]].corr()['Class'][:-1]
plt.plot(corr)
plt.xticks(rotation=90)
plt.show()
```





```
[102]: X = data_clean2.drop('Class', axis=1)
       y = data_clean2['Class']
```

```
[103]: from sklearn.model_selection import train_test_split
       from sklearn.metrics import accuracy_score
       X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.
       ↪2,random_state=0)
```

```
[104]: from sklearn.ensemble import RandomForestClassifier
       rfc = RandomForestClassifier(random_state=0)
       rfc.fit(X_train, y_train)
```

```
[104]: RandomForestClassifier(random_state=0)
```

```
[105]: y_pred = rfc.predict(X_test)
       print("Accuracy Score :", round(accuracy_score(y_test, y_pred)*100 ,2), "%")
```


Accuracy Score : 100.0 %

```
[106]: from sklearn.metrics import accuracy_score, f1_score, precision_score, recall_score
        print('F-1 Score : ',(f1_score(y_test, y_pred)))
        print('Precision Score : ',(precision_score(y_test, y_pred)))
        print('Recall Score : ',(recall_score(y_test, y_pred)))
```

```
F-1 Score : 1.0
Precision Score : 1.0
Recall Score : 1.0
```

```
[107]: from sklearn.metrics import classification_report, confusion_matrix
        cm = confusion_matrix(y_test, y_pred)
        plt.figure(figsize=(5,5))
        sns.heatmap(data=cm,linewidths=.5, annot=True,square = True, cmap = 'Blues')
        plt.ylabel('Actual label')
        plt.xlabel('Predicted label')
        all_sample_title = 'Accuracy Score: {0}'.format(rfc.score(X_test, y_test)*100)
        plt.title(all_sample_title, size = 15)
```

```
[107]: Text(0.5, 1.0, 'Accuracy Score: 100.0')
```



```
[108]: from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier()
knn.fit(X_train, y_train)
```

```
[108]: KNeighborsClassifier()
```

```
[109]: y_pred = knn.predict(X_test)
print("Accuracy Score :", round(accuracy_score(y_test, y_pred)*100 ,2), "%")
```


Accuracy Score : 83.33 %

```
[110]: from sklearn.metrics import accuracy_score, f1_score, precision_score, r
↳ recall_score
print('F-1 Score : ',(f1_score(y_test, y_pred)))
print('Precision Score : ',(precision_score(y_test, y_pred)))
print('Recall Score : ',(recall_score(y_test, y_pred)))
```

F-1 Score : 0.7586206896551724  
Precision Score : 0.88  
Recall Score : 0.6666666666666666

```
[111]: from sklearn.metrics import classification_report, confusion_matrix
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(5,5))
sns.heatmap(data=cm,linewidths=.5, annot=True,square = True, cmap = 'Blues')
plt.ylabel('Actual label')
plt.xlabel('Predicted label')
all_sample_title = 'Accuracy Score: {0}'.format(knn.score(X_test, y_test)*100)
plt.title(all_sample_title, size = 15)
```

```
[111]: Text(0.5, 1.0, 'Accuracy Score: 83.33333333333334')
```



```
[112]: from sklearn.ensemble import AdaBoostClassifier
ada = AdaBoostClassifier(random_state=0)
ada.fit(X_train, y_train)
```

C:\Users\dell\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\ensemble\\_weight\_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.

```
warnings.warn(
```

```
[112]: AdaBoostClassifier(random_state=0)
```

```
[113]: y_pred = ada.predict(X_test)
print("Accuracy Score :", round(accuracy_score(y_test, y_pred)*100 ,2), "%")
```


Accuracy Score : 100.0 %

```
[114]: from sklearn.metrics import accuracy_score, f1_score, precision_score, r
        ↪ recall_score
print('F-1 Score : ',(f1_score(y_test, y_pred)))
print('Precision Score : ',(precision_score(y_test, y_pred)))
print('Recall Score : ',(recall_score(y_test, y_pred)))
```

F-1 Score : 1.0  
Precision Score : 1.0  
Recall Score : 1.0

```
[115]: from sklearn.metrics import classification_report, confusion_matrix
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(5,5))
sns.heatmap(data=cm,linewidths=.5, annot=True,square = True, cmap = 'Blues')
plt.ylabel('Actual label')
plt.xlabel('Predicted label')
all_sample_title = 'Accuracy Score: {0}'.format(ada.score(X_test, y_test)*100)
plt.title(all_sample_title, size = 15)
```

[115]: Text(0.5, 1.0, 'Accuracy Score: 100.0')



```
[116]: from sklearn.linear_model import LogisticRegression
lr = LogisticRegression(random_state = 0)
lr.fit(X_train, y_train)
```

C:\Users\dell\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\linear\_model\\_logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
[116]: LogisticRegression(random_state=0)
```

```
[117]: y_pred = lr.predict(X_test)
print("Accuracy Score :", round(accuracy_score(y_test, y_pred)*100 ,2), "%")
```

Accuracy Score : 100.0 %

```
[118]: from sklearn.metrics import accuracy_score, f1_score, precision_score, r
        ↪recall_score
print('F-1 Score : ',(f1_score(y_test, y_pred)))
print('Precision Score : ',(precision_score(y_test, y_pred)))
print('Recall Score : ',(recall_score(y_test, y_pred)))
```


F-1 Score : 1.0

Precision Score : 1.0

Recall Score : 1.0

```
[119]: from sklearn.metrics import classification_report, confusion_matrix
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(5,5))
sns.heatmap(data=cm,linewidths=.5, annot=True,square = True, cmap = 'Blues')
plt.ylabel('Actual label')
plt.xlabel('Predicted label')
all_sample_title = 'Accuracy Score: {0}'.format(lr.score(X_test, y_test)*100)
plt.title(all_sample_title, size = 15)
```

```
[119]: Text(0.5, 1.0, 'Accuracy Score: 100.0')
```



[ ]: